

Assignment 2 - MAD

Mikkel Willén

November 2021

Exercise 1

a)

Vi har

$$A = (\mathbf{a} \cdot \mathbf{1}^T) \cdot \mathbf{I}, \quad \mathbf{a} = [a_1, a_2, \dots, a_n]^T, \quad \mathbf{1} = [1, 1, \dots, 1]^T$$

Vi kan derfor skrive funktionen som

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} A(X\mathbf{w} - \mathbf{t})^T (X\mathbf{w} - \mathbf{t})$$

Vi simplificere funktionen

$$\begin{aligned} \mathcal{L}(\mathbf{w}) &= \frac{1}{N} A(X\mathbf{w} - \mathbf{t})^T (X\mathbf{w} - \mathbf{t}) \\ &= \frac{1}{N} A((X\mathbf{w})^T - \mathbf{t}^T)(X\mathbf{w} - \mathbf{t}) \\ &= \frac{1}{N} A(X\mathbf{w})^T X\mathbf{w} - \frac{1}{N} \mathbf{t}^T X\mathbf{w} - \frac{1}{N} A(X\mathbf{w})^T \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{t} \\ &= \frac{1}{N} A\mathbf{w}^T X^T X\mathbf{w} - \frac{2}{N} A\mathbf{w}^T X^T \mathbf{t} + \frac{1}{N} A\mathbf{t}^T \mathbf{t} \end{aligned}$$


Vi finder nu gradienten for denne funktion

$$\begin{aligned} \nabla \mathcal{L}(\mathbf{w}) &= \frac{2}{N} AX^T X\mathbf{w} - \frac{2}{N} AX^T \mathbf{t} \\ \nabla \mathcal{L}(\mathbf{w}) &= 0 \\ \Leftrightarrow \quad \frac{2}{N} AX^T X\mathbf{w} - \frac{2}{N} AX^T \mathbf{t} &= 0 \\ \Leftrightarrow \quad AX^T X\mathbf{w} &= AX^T \mathbf{t} \\ \Leftrightarrow \quad \mathbf{I}\mathbf{w} &= (X^T AX)^{-1} X^T A\mathbf{t} \\ &\quad \hat{\mathbf{w}} = (X^T AX)^{-1} X^T A\mathbf{t} \end{aligned}$$

b)

Vi forventede, at vi ville få en bedre løsning, end i den forrige opgave, men det har vi ikke fået. RMSE'en for den nye løsning er større end i den forrige. Disse vægte ændre markant på, hvad vi får ud af funktionerne.


Koden kan ses i appendix. Herunder er et billede, af vores scatterplot



Exercise 2

Se kode i appendix.

a)



Bedste lambdaværdi: 0.0000003430

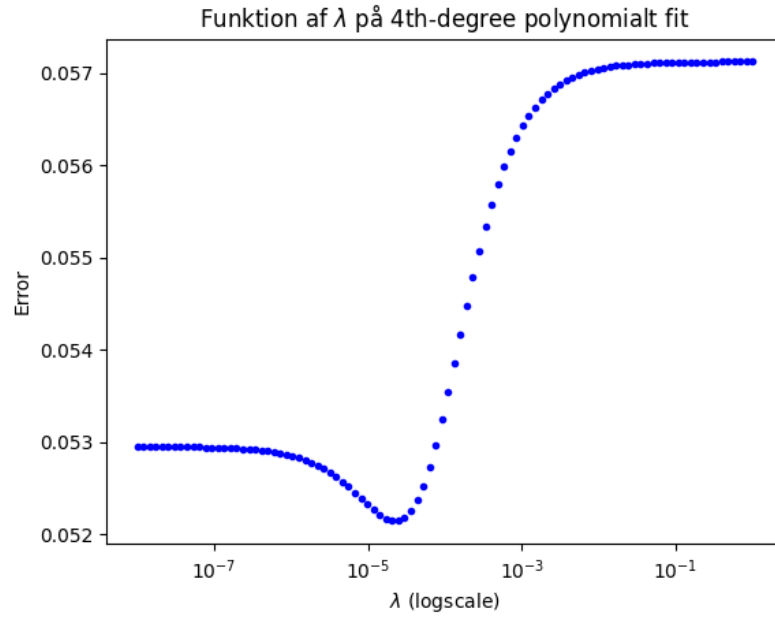
Værdier for $\lambda = 0$

$$\begin{bmatrix} 3.64164559 \cdot 10 \\ -1.33308857 \cdot 10^{-2} \end{bmatrix}$$

Værdier for $\lambda = 0.0000003430$

$$\begin{bmatrix} 3.63776282 \cdot 10 \\ -1.33110046e \cdot 10^{-2} \end{bmatrix}$$

b)



Bedste lambdaværdi: 0.0000205651

Værdier for $\lambda = 0$

$$\begin{bmatrix} 3.21358282 \cdot 10^5 \\ -6.46374641 \cdot 10^2 \\ 4.87449415 \cdot 10^{-1} \\ -1.63339180 \cdot 10^{-4} \\ 2.05197882 \cdot 10^{-8} \end{bmatrix}$$

Værdier for $\lambda = 0.0000205651$

$$\begin{bmatrix} 1.88300350 \cdot 10^{-2} \\ 9.11277821 \\ -1.38600662 \cdot 10^{-2} \\ 7.03376807e \cdot 10^{-6} \\ -1.19035006 \cdot 10^{-9} \end{bmatrix}$$

Exercise 3

a)

Vi differentiere funktionen

$$f(x) = 1 - e^{-\beta x^\alpha}$$

$$\begin{aligned} \frac{d}{dx} (1 - e^{-\beta x^\alpha}) &= -\frac{d}{dx} (e^{-\beta x^\alpha}) && \text{konstant regl og sumregl} \\ &= \frac{d}{du} e^u \cdot u', \quad u = -\beta x^\alpha && \text{kæderegl} \\ &= e^u \cdot u', \quad u' = \left(-\frac{\beta x^\alpha \alpha}{x}\right) && \text{diff af begge funktioner} \\ &= e^{-\beta x^\alpha} \cdot \left(-\frac{\beta x^\alpha \alpha}{x}\right) && \text{sætter ind} \\ &= e^{-\beta x^\alpha} \beta \alpha x^{\alpha-1} && \text{reduceret} \end{aligned}$$

Skriver pdf'en op

$$F(x) = \begin{cases} 0 & x \leq 0 \\ e^{-\beta x^\alpha} \beta \alpha x^{\alpha-1} & x > 0 \end{cases}$$

b)

Vi sætter $\beta = \frac{1}{4}$ og $\alpha = 2$

$$F(x) = e^{-\frac{1}{4}x^2} \frac{1}{2}x$$

og udregner $F(4)$

$$F(4) = e^{-\frac{1}{4}4^2} \frac{1}{2}4 = 2 \cdot e^{-4} = 0.0366$$

herefter udregner vi $F(5) - F(10)$

$$F(5) - F(10) = e^{-\frac{1}{4}5^2} \frac{1}{2}5 - \left(e^{-\frac{1}{4}10^2} \frac{1}{2}10\right) = 0.00483$$

c)

Vi finder integralet for pdf'en, som er cdf'en. Så løser vi nedenstående ligning for m

$$\begin{aligned} 1 - e^{-\beta m^\alpha} - \left(1 - e^{-\beta 0^\alpha}\right) &= \frac{1}{2} \\ \Leftrightarrow 1 - e^{-\beta m^\alpha} &= \frac{1}{2} \\ \Leftrightarrow m &= e^{\frac{\ln\left(\frac{\ln(2)}{\beta}\right)}{\alpha}} \end{aligned}$$

Exercise 4

a)

I denne opgave antager jeg, at der ikke skulle stå "not" i opgavebeskrivelsen da "not speak" og "silent" virker som det samme.

Jeg ganger de forskellige faktorer sammen og får:

$$X_{speak} = 0.002 \cdot 0.5 \cdot 0.75 = 0.00075$$

b)

Jeg ganger de forskellige faktorer sammen og får:

$$X_{silent} = 0.001 \cdot (1 - 0.5 \cdot 0.25) = 0.000875$$

Det giver derfor mening at sige noget, da det giver en mindre potentiel straf.

c)

jeg ganger de forskellige faktorer sammen og får:

$$Y_{speak} = 0.005 \cdot (1 - 0.1) \cdot 0.75 = 0.003375$$

$$Y_{silent} = 0.001 \cdot (1 - 0.1 \cdot 0.25) = 0.000975$$

Det giver derfor mening ikke at sige noget, da det giver en mindre potentiel straf.

Appendix

linweightreg.py

```
import numpy

class LinearRegression():
    """
    Linear regression implementation.
    """

    def __init__(self):

        pass

    def fit(self, X, t):
        """
        Fits the linear regression model.

        Parameters
        -----
        X : Array of shape [n_samples, n_features]
        t : Array of shape [n_samples, 1]
        """
        A = numpy.zeros((X.shape[0], X.shape[0]))
        for i in range(X.shape[0]):
            A[i][i] = t[i] ** 2
        print(A)
        X = numpy.reshape(X, (X.shape[0], -1))
        X = numpy.insert(X, 0, 1, axis=1)
        XTX = numpy.dot(numpy.dot(X.T, A), X)
        inverse = numpy.linalg.inv(XTX)
        XTXX = numpy.dot(inverse, X.T)
        self.w = numpy.dot(numpy.dot(XTXX, A), t)

    def predict(self, X):
        """
        Computes predictions for a new set of points.

        Parameters
        -----
        X : Array of shape [n_samples, n_features]
```

```

Returns
-----
predictions : Array of shape [n_samples, 1]
"""
X = numpy.reshape(X, (X.shape[0], -1))
X = numpy.insert(X, 0, 1, axis = 1)
self.p = numpy.dot(X, self.w)

```

housing_2.py

```

import numpy
import pandas
import linweightreg
import matplotlib.pyplot as plt

# load data
train_data = numpy.loadtxt("boston_train.csv", delimiter=",")
test_data = numpy.loadtxt("boston_test.csv", delimiter=",")
X_train, t_train = train_data[:, :-1], train_data[:, -1]
X_test, t_test = test_data[:, :-1], test_data[:, -1]
# make sure that we have N-dimensional Numpy arrays (ndarray)
t_train = t_train.reshape((len(t_train), 1))
t_test = t_test.reshape((len(t_test), 1))
print("Number of training instances: %i" % X_train.shape[0])
print("Number of test instances: %i" % X_test.shape[0])
print("Number of features: %i" % X_train.shape[1])

# (b) fit linear regression using only the first feature
model_single = linweightreg.LinearRegression()
model_single.fit(X_train[:, 0], t_train)
print("weights for a single variable: \n", model_single.w)
print("Husprisen starter på 23 tusinde dollars, når crimerate er 0, og falder med 0.4,
      når enheden for crimerate stiger med 1")

# (c) fit linear regression model using all features
model_multiple = linweightreg.LinearRegression()
model_multiple.fit(X_train, t_train)
print("weights for multiple variables: \n", model_multiple.w)
print("Husprisen starter på 31 tusiden dollars, når alle andre værdier er 0,
      og falder eller stiger, alt efter fortegnet på værdien, når de andre værdier stiger")

# (d) evaluation of results
model_single.predict(X_test[:, 0])
model_multiple.predict(X_test)

def rmse(t, tp):

```

```

N = t.shape[0]
sum = 0.0
for i in range (N):
    sum = sum + (t[i] - tp[i]) ** 2
rmse = numpy.sqrt(sum/N)
return rmse

print("RMSE for a single variable: ", rmse(t_test, model_single.p))
print("RMSE for multiple variables: ", rmse(t_test, model_multiple.p))

print("Plot of single variable")
plt.scatter(model_single.p, t_test)
print("Plot of multiple variables")
plt.scatter(model_multiple.p, t_test)

```

linreg.py

```

import numpy

class LinearRegression():
    """
    Linear regression implementation with
    regularization.
    """

    def __init__(self, lam=0, solver="inverse"):
        self.lam = lam
        self.solver = solver

        assert self.solver in ["inverse", "solve"]

    def fit(self, X, t):
        """
        Fits the linear regression model.

        Parameters
        -----
        X : Array of shape [n_samples, n_features]
        t : Array of shape [n_samples, 1]
        """

        X = numpy.reshape(X, (X.shape[0], -1))
        X = numpy.insert(X, 0, 1, axis=1)
        t = numpy.reshape(t, (len(t),1))

        diagonal = self.lam * len(X) * numpy.identity(X.shape[1])

```



```

km = numpy.dot(X.T, X) + diagonal

if self.solver == "solve":
    self.w = numpy.linalg.solve(km, numpy.dot(X.T, t))

elif self.solver == "inverse":
    self.w = numpy.linalg.inv(km)
    self.w = numpy.dot(self.w, X.T)
    self.w = numpy.dot(self.w, t)

else:
    raise Exception("Unknown solver!")

def predict(self, X):
    """
    Computes predictions for a new set of points.

    Parameters
    -----
    X : Array of shape [n_samples, n_features]

    Returns
    -----
    predictions : Array of shape [n_samples, 1]
    """
    X = numpy.reshape(X, (X.shape[0], -1))
    X = numpy.insert(X, 0, 1, axis = 1)
    self.p = numpy.dot(X, self.w)

```

RegularizedLinReg.py

```

import matplotlib.pyplot as plt
import numpy as np
import linreg

solver = "solve"

def lossFunction(X, t, lam, verbose=0):

    loss = 0

    for i in range(len(X)):
        X_train = np.delete(X, i, 0)
        t_train = np.delete(t, i, 0)

        model = linreg.LinearRegression(lam=lam, solver=solver)

```

```

        model.fit(X_train, t_train)

        X_val = X[i].reshape((1, X_train.shape[1]))
        t_val = t[i].reshape((1, 1))
        model.predict(X_val)

        loss += (t_val[0,0] - model.p[0,0]) ** 2.0

    loss = loss / len(X)

    if verbose > 0:
        print("lam=%.10f and loss=%.10f" % (lam, loss))

    return loss

raw = np.genfromtxt('men-olympics-100.txt', delimiter=' ')

t = raw[:,1].reshape((len(raw),1))

lambdaValues = np.logspace(-8, 0, 100, base=10)
print("1st Order Polynomial: ")
X = raw[:,0].reshape((len(raw),1))

resultA = np.array([lossFunction(X, t, lam) for lam in lambdaValues])
bestLambdaA = lambdaValues[np.argmin(resultA)]
print("Best lambda value: %.10f" % bestLambdaA)

modelZero = linreg.LinearRegression(lam=0.0, solver=solver)
modelZero.fit(X, t)
print("Optimal coefficients for lam=%.10f: \n%s" % (0.0, str(modelZero.w)))

modelBest = linreg.LinearRegression(lam=bestLambdaA, solver=solver)
modelBest.fit(X, t)
print("Optimal coefficients for lam=%.10f: \n%s" % (bestLambdaA, str(modelBest.w)))

plt.figure()
plt.plot(lambdaValues, resultA, "bo", markersize=3)
plt.title("Funktion af  $\lambda$  på 1st-degree polynomial fit")
plt.ylabel("Error")
plt.xlabel(" $\lambda$  (logscale)")
plt.xscale("log")
plt.show()

print("4th Degree Polynomial")
X4 = np.empty((len(raw[:,0]),4))

```

```

X4[:,0] = raw[:,0]
X4[:,1] = raw[:,0] ** 2
X4[:,2] = raw[:,0] ** 3
X4[:,3] = raw[:,0] ** 4

resultB = np.array([lossFunction(X4, t, lam) for lam in lambdaValues])
bestLambdab = lambdaValues[np.argmin(resultB)]
print("Best lambda value: %.10f" % bestLambdab)

modelZero = linreg.LinearRegression(lam=0.0, solver=solver)
modelZero.fit(X4, t)
print("Optimal coefficients for lam=%.10f: \n%s" % (0.0, str(modelZero.w)))

modelBest = linreg.LinearRegression(lam=bestLambdab, solver=solver)
modelBest.fit(X4, t)
print("Optimal coefficients for lam=%.10f: \n%s" % (bestLambdab, str(modelBest.w)))

plt.figure()
plt.plot(lambdaValues, resultB, "bo", markersize=3)
plt.title("Funktion af  $\lambda$  på 4th-degree polynomial fit")
plt.ylabel("Error")
plt.xlabel(" $\lambda$  (logscale)" )
plt.xscale("log")
plt.show()

```